



Cover Memo: Firm AI Use Policy Package

To: Suzanne Babb

From: John A. Kelly

Date: August 5, 2026

Re: Model AI use policy and court disclosure package for your firm (research assistance)

Dear Suzanne,

I put together, as a research project, a package your firm can work from and adapt into its own AI use policy, with a court disclosure component. This is research assistance from one practitioner to another, built from public authorities and my own work on California AI ethics and local rules. I think we both agree this is not legal advice to your firm; the bracketed decisions and final adoption belong to your firm and its counsel.

Why now

Two independent drivers. First, the State Bar's COPRAC issued its updated Practical Guidance in 2026, and it treats firm AI policies as a managerial duty: under Rules 5.1 and 5.3, supervising lawyers "should establish clear policies regarding the uses of generative AI," train staff, and keep the policy current, with heightened duties for agentic (autonomous) tools. A firm without a written policy is now visibly behind the standard of care; research indicates written gen-AI policies exist at fewer than half of firms, so adopting one now is also a differentiator. Second, county local rules have arrived, including six California counties that have adopted AI provisions effective in 2026, and Santa Clara now mandates disclosure of AI use in ex parte filings. SB 574, pending in the Assembly, would put a personal read-and-verify duty into Code of Civil Procedure section 128.7 statewide.

What is in the package

1. **Model Firm AI Use Policy** (document 02). Twelve sections mapped to the COPRAC duty structure: governance and an approved tool list, confidentiality input tiers, memory and cross-matter contamination controls, rules for AI add-ins working inside client documents, the four-step citation verification ladder, supervision and training, client communication, billing, court compliance, documentation, bias, and periodic review. Bracketed items mark the firm's decisions.
2. **Court Disclosure Kit** (document 03). Three instruments: a Case Management Statement attachment (CM-110) making categorical disclosures once per case; an optional per-filing AI Disclosure Certificate; and a Santa Clara ex parte compliance note. The design premise is that disclosure at the CMC names the accountable human on day one, the way the profession absorbed e-discovery, while the section 128.7 signature continues to govern every individual filing.
3. **County Compliance Matrix** (document 04). The six counties with AI local rules, verified against the July 2026 rule texts, plus the statewide layer and pending developments.
4. **Staff Quick Reference** (document 05). A one-page version for daily use.
5. **Placeholder Card** (document 07). The practical answer to "strip client identifiers": the standard substitution conventions (Doe and Roe parties, Blackacre, Acme Corp.), the rule that substitution happens before text reaches any cloud tool, and the account posture that makes Tier 2 work (enterprise configuration, training off, zero data retention where offered, and knowing your vendor's actual retention window).
6. **Verification Prompt Pack** (document 08). Copy-and-paste prompts for both failure modes: legal (the four-

step citation ladder, closed-corpus research) and factual (a two-session claim-inventory check of any draft against its own source documents). Everything runs in an ordinary chat interface, including the Claude add-in in Word; nothing requires anyone to touch code.

7. **Grounding Sources and Free Case Retrieval** (document 09). Where to get authoritative text to ground the tools (maintained plain-text repositories of the statutes, the Rules of Court, and all 58 counties' local rules, plus the official sources), and the free retrieval chain for cases outside a research plan, which this package's own cite-check used to pull two federal sanctions opinions at no charge.
8. **Adoption Notes for Firm Management** (document 10). The market context that goes stale, kept out of the policy so the policy survives price changes: cyber insurance expectations, budget benchmarks, seat-versus-token pricing and its deadline risk, and a training cadence model.
9. **Where the Data Lives** (document 11). The three ways firms actually work with AI (web chat, Word add-ins, agentic harnesses) analyzed for actual data residence: what transits to the vendor, where conversations and memory persist, where zero data retention is and is not achievable, and which policy controls carry the load in each modality. The short version: the add-in that looks the most ordinary has the hardest confidentiality profile, and the harness that looks the most exotic has the most controllable one.

Added since this memo was first drafted (updated August 11, 2026)

The package has grown three companion sets, all organized on the materials page that accompanies this memo:

1. **A presentation deck** (24 slides, plain-language edition) built to walk a partnership through the package in about an hour, with a standalone glossary handout that defines every term of art the materials use.
2. **Product policy documents**: eleven per-product implementations of the model policy (ChatGPT Business, Claude Team, Microsoft 365 Copilot, Cursor, embedded AI features, CoCounsel, Lexis+ AI Protege, Vincent, Harvey, Legora), each showing exactly how the policy lands on that tool: the setup steps, what may go in, and the vendor's actual terms, graded by whether they bind.
3. **A one-page cost and configuration sheet**: what each product costs per seat, the commitment floors, and how long each takes to stand up, with reported (non-published) prices clearly flagged.

Next steps for tailoring:

1. Given your firm size and structure, who would be the AI Policy Administrator.
2. Practice areas and the counties where you regularly file. I'd presume Sonoma, Napa, Marin and Mendocino, but good to have the list you'd want to work from.
3. What AI tools people are actually using today, and on what accounts (enterprise or personal). The honest inventory matters more than the aspirational one.
4. Whether any clients impose AI restrictions in engagement terms or outside counsel guidelines.
5. Your appetite on voluntary disclosure: default to the CMC attachment in every litigated matter, or reserve it for matters where AI use is material? The kit works either way; the policy should state a default.

Watch items

The COPRAC proposed amendments to Rules 1.1, 1.4, 1.6, and 3.3 (comment period closed in May; Trustees action expected later this year) and SB 574 (Assembly deadline August 31, 2026) could each turn parts of this from best practice into binding obligation. The package is written so that neither would require structural change, only a citation update.

Happy to walk through it whenever suits. Appreciate you and Dave both for your continuing participation in all this

AI stuff.

Sincerely,

John